

```
In [1]: import numpy as np # linear algebra, matrix operations
import pandas as pd # to read data
import matplotlib.pyplot as plt # plotting, showing images
```

```
In [2]: from keras.datasets import mnist
```

```
In [3]: (X_train, Y_train), (X_test, Y_test) = mnist.load_data() # import data
print(X_train.shape, Y_train.shape, X_test.shape, Y_test.shape)

(60000, 28, 28) (60000,) (10000, 28, 28) (10000,)
```

```
In [4]: X_train = X_train.flatten().reshape(X_train.shape[0], X_train.shape[1] * X_train.shape[2]) # convert to 2D
X_test = X_test.flatten().reshape(X_test.shape[0], X_test.shape[1] * X_test.shape[2]) # convert to 2D
Y_train = Y_train.reshape(Y_train.shape[0], 1) # convert to 2D
Y_test = Y_test.reshape(Y_test.shape[0], 1) # convert to 2D
print(X_train.shape, Y_train.shape, X_test.shape, Y_test.shape)

(60000, 784) (60000, 1) (10000, 784) (10000, 1)
```

```
In [5]: X_train = np.hstack((Y_train, X_train)) # stack data horizontally
X_test = np.hstack((Y_test, X_test)) # stack data horizontally
print(X_train.shape, X_test.shape)

(60000, 785) (10000, 785)
```

```
In [6]: m, n = X_train.shape # n,m inverted because we will transpose
np.random.shuffle(X_train) # shuffle for randomness

dev_data = X_train[0:5000].T # dev data used for hyperparam tuning
Y_dev = dev_data[0]
X_dev = dev_data[1:n]
X_dev = X_dev / 255.0

train_data = X_train[5000:].T
Y_train = train_data[0]
X_train = train_data[1:n]
X_train = X_train / 255.0

n_train, m_train = X_train.shape
```

```
In [7]: def init_params():
    W1 = np.random.rand(10, 784) - 0.5
    b1 = np.random.rand(10, 1) - 0.5
    W2 = np.random.rand(10, 10) - 0.5
    b2 = np.random.rand(10, 1) - 0.5
    return W1, b1, W2, b2

def ReLU(Z):
    return np.maximum(Z, 0)

def softmax(Z):
    A = np.exp(Z) / sum(np.exp(Z))
    return A

def forward_prop(W1, b1, W2, b2, X):
    Z1 = W1.dot(X) + b1
    A1 = ReLU(Z1)
    Z2 = W2.dot(A1) + b2
    A2 = softmax(Z2)
    return Z1, A1, Z2, A2

def ReLU_deriv(Z):
    return Z > 0

def one_hot(Y):
    one_hot_Y = np.zeros((Y.size, Y.max() + 1))
    one_hot_Y[np.arange(Y.size), Y] = 1
    one_hot_Y = one_hot_Y.T
    return one_hot_Y
```

```

def backward_prop(Z1, A1, Z2, A2, W1, W2, X, Y):
    one_hot_Y = one_hot(Y)
    dZ2 = A2 - one_hot_Y
    dW2 = 1 / m * dZ2.dot(A1.T)
    db2 = 1 / m * np.sum(dZ2)
    dZ1 = W2.T.dot(dZ2) * ReLU_deriv(Z1)
    dW1 = 1 / m * dZ1.dot(X.T)
    db1 = 1 / m * np.sum(dZ1)
    return dW1, db1, dW2, db2

def update_params(W1, b1, W2, b2, dW1, db1, dW2, db2, alpha):
    W1 = W1 - alpha * dW1
    b1 = b1 - alpha * db1
    W2 = W2 - alpha * dW2
    b2 = b2 - alpha * db2
    return W1, b1, W2, b2

```

```

In [8]: def get_predictions(A2):
        return np.argmax(A2, 0)

def get_accuracy(predictions, Y):
    print(predictions, Y)
    return np.sum(predictions == Y) / Y.size

def gradient_descent(X, Y, alpha, iterations):
    W1, b1, W2, b2 = init_params()
    for i in range(iterations):
        Z1, A1, Z2, A2 = forward_prop(W1, b1, W2, b2, X)
        dW1, db1, dW2, db2 = backward_prop(Z1, A1, Z2, A2, W1, W2, X, Y)
        W1, b1, W2, b2 = update_params(W1, b1, W2, b2, dW1, db1, dW2, db2, alpha)
        if i % 10 == 0:
            print("Iteration: ", i)
            predictions = get_predictions(A2)
            print(get_accuracy(predictions, Y))
    return W1, b1, W2, b2

```

```

In [10]: W1, b1, W2, b2 = gradient_descent(X_train, Y_train, 0.10, 500)

```

Iteration: 0  
[0 0 0 ... 3 5 5] [6 5 7 ... 9 0 2]  
0.09989090909090909  
Iteration: 10  
[0 0 0 ... 3 5 3] [6 5 7 ... 9 0 2]  
0.17103636363636363  
Iteration: 20  
[0 0 0 ... 3 5 8] [6 5 7 ... 9 0 2]  
0.24354545454545454  
Iteration: 30  
[0 0 0 ... 3 5 7] [6 5 7 ... 9 0 2]  
0.31616363636363637  
Iteration: 40  
[0 0 0 ... 3 5 7] [6 5 7 ... 9 0 2]  
0.38127272727272726  
Iteration: 50  
[0 0 7 ... 3 5 7] [6 5 7 ... 9 0 2]  
0.4339636363636364  
Iteration: 60  
[0 0 7 ... 3 5 7] [6 5 7 ... 9 0 2]  
0.48165454545454545  
Iteration: 70  
[0 0 7 ... 3 5 7] [6 5 7 ... 9 0 2]  
0.5237090909090909  
Iteration: 80  
[0 0 7 ... 3 5 8] [6 5 7 ... 9 0 2]  
0.5606727272727273  
Iteration: 90  
[0 0 7 ... 3 2 2] [6 5 7 ... 9 0 2]  
0.5942  
Iteration: 100  
[2 0 7 ... 3 2 2] [6 5 7 ... 9 0 2]  
0.6237272727272727  
Iteration: 110  
[2 0 7 ... 3 2 2] [6 5 7 ... 9 0 2]  
0.6478  
Iteration: 120  
[2 0 7 ... 9 2 2] [6 5 7 ... 9 0 2]  
0.6693818181818182  
Iteration: 130  
[2 0 7 ... 9 2 2] [6 5 7 ... 9 0 2]  
0.6877272727272727  
Iteration: 140  
[2 0 7 ... 9 2 2] [6 5 7 ... 9 0 2]  
0.7038  
Iteration: 150  
[2 0 7 ... 9 2 2] [6 5 7 ... 9 0 2]  
0.718  
Iteration: 160  
[2 0 7 ... 9 2 2] [6 5 7 ... 9 0 2]  
0.7296727272727273  
Iteration: 170  
[2 0 7 ... 9 2 2] [6 5 7 ... 9 0 2]  
0.7407636363636364  
Iteration: 180  
[2 0 7 ... 9 2 2] [6 5 7 ... 9 0 2]  
0.7495818181818181  
Iteration: 190  
[2 0 7 ... 9 2 2] [6 5 7 ... 9 0 2]  
0.7569818181818182  
Iteration: 200  
[2 0 7 ... 9 2 2] [6 5 7 ... 9 0 2]  
0.7653272727272727  
Iteration: 210  
[2 0 7 ... 9 2 2] [6 5 7 ... 9 0 2]  
0.7715454545454545  
Iteration: 220  
[2 0 7 ... 9 2 2] [6 5 7 ... 9 0 2]  
0.7773272727272728  
Iteration: 230

[2 0 7 ... 9 2 2] [6 5 7 ... 9 0 2]  
0.7827090909090909  
Iteration: 240  
[2 0 7 ... 9 2 2] [6 5 7 ... 9 0 2]  
0.7875636363636364  
Iteration: 250  
[2 0 7 ... 9 2 2] [6 5 7 ... 9 0 2]  
0.7917454545454545  
Iteration: 260  
[2 0 7 ... 9 2 2] [6 5 7 ... 9 0 2]  
0.7964909090909091  
Iteration: 270  
[6 0 7 ... 9 2 2] [6 5 7 ... 9 0 2]  
0.8005272727272728  
Iteration: 280  
[6 0 7 ... 9 2 2] [6 5 7 ... 9 0 2]  
0.8042363636363636  
Iteration: 290  
[6 0 7 ... 9 2 2] [6 5 7 ... 9 0 2]  
0.8076727272727273  
Iteration: 300  
[6 0 7 ... 9 2 2] [6 5 7 ... 9 0 2]  
0.8105454545454546  
Iteration: 310  
[6 0 7 ... 9 2 2] [6 5 7 ... 9 0 2]  
0.8132909090909091  
Iteration: 320  
[6 0 7 ... 9 2 2] [6 5 7 ... 9 0 2]  
0.8162545454545455  
Iteration: 330  
[6 0 7 ... 9 2 2] [6 5 7 ... 9 0 2]  
0.8190181818181819  
Iteration: 340  
[6 0 7 ... 9 2 2] [6 5 7 ... 9 0 2]  
0.8218  
Iteration: 350  
[6 0 7 ... 9 2 2] [6 5 7 ... 9 0 2]  
0.824  
Iteration: 360  
[6 0 7 ... 9 2 2] [6 5 7 ... 9 0 2]  
0.8261272727272727  
Iteration: 370  
[6 0 7 ... 9 2 2] [6 5 7 ... 9 0 2]  
0.8285090909090909  
Iteration: 380  
[6 0 7 ... 9 2 2] [6 5 7 ... 9 0 2]  
0.8302727272727273  
Iteration: 390  
[6 0 7 ... 9 2 2] [6 5 7 ... 9 0 2]  
0.8315818181818182  
Iteration: 400  
[6 0 7 ... 9 2 2] [6 5 7 ... 9 0 2]  
0.8335636363636364  
Iteration: 410  
[6 0 7 ... 9 2 2] [6 5 7 ... 9 0 2]  
0.8351272727272727  
Iteration: 420  
[6 0 7 ... 9 2 2] [6 5 7 ... 9 0 2]  
0.8368909090909091  
Iteration: 430  
[6 0 7 ... 9 2 2] [6 5 7 ... 9 0 2]  
0.8386  
Iteration: 440  
[6 0 7 ... 9 2 2] [6 5 7 ... 9 0 2]  
0.8402727272727273  
Iteration: 450  
[6 0 7 ... 9 2 2] [6 5 7 ... 9 0 2]  
0.8416  
Iteration: 460  
[6 0 7 ... 9 2 2] [6 5 7 ... 9 0 2]

```

0.8431272727272727
Iteration: 470
[6 0 7 ... 9 2 2] [6 5 7 ... 9 0 2]
0.8442545454545455
Iteration: 480
[6 0 7 ... 9 8 2] [6 5 7 ... 9 0 2]
0.8454
Iteration: 490
[6 0 7 ... 9 8 2] [6 5 7 ... 9 0 2]
0.8463454545454545

```

```

In [11]: def make_predictions(X, W1, b1, W2, b2):
_, _, A2 = forward_prop(W1, b1, W2, b2, X)
predictions = get_predictions(A2)
return predictions

def test_prediction(index, W1, b1, W2, b2):
current_image = X_train[:, index, None]
prediction = make_predictions(X_train[:, index, None], W1, b1, W2, b2)
label = Y_train[index]
print("Prediction: ", prediction)
print("Label: ", label)

current_image = current_image.reshape((28, 28)) * 255
plt.gray()
plt.imshow(current_image, interpolation='nearest')
plt.show()

```

```


In [12]: test_prediction(0, W1, b1, W2, b2)
test_prediction(1, W1, b1, W2, b2)
test_prediction(2, W1, b1, W2, b2)
test_prediction(3, W1, b1, W2, b2)

```

```

Prediction: [6]
Label: 6


```



```


Prediction: [0]
Label: 5

```




Prediction: [7]

Label: 7



Prediction: [8]

Label: 8



```
In [13]: dev_predictions = make_predictions(X_dev, W1, b1, W2, b2)
         get_accuracy(dev_predictions, Y_dev)
```

```
[9 9 7 ... 7 2 6] [4 9 0 ... 7 2 6]
```

```
Out[13]: 0.8456
```

```
In [ ]:
```